



# Agenda

- Intro to Delta Lake

# Intro to Delta Lake

# What are traditional Data Lakes?

- They are storage repositories that stores a large amount of raw data(current and historical) in native formats
- May also contain relational databases with live transactional data.
- Versatile, scalable, cheap.

# Problems with traditional Data Lakes

- Not ready for direct data analysis and ML
- Hard to enforce schema for data, prone to inconsistent and low quality data.
- Difficulty sorting data by an index if data is spread across many files and partitioned.
- Failed jobs might leave data in corrupt states
- Too much overhead opening & reading when working with large amount of small files
- Partitioning is only “poor man’s indexing”
- No caching, low throughput

# What is Delta Lake?

- **Data Lake + Delta file format = Delta Lake**
- A file format that's designed specifically to work with Spark and DBFS
- Has both open-source and managed offerings
- Data is stored as **Parquet** files, and a **transaction log** is maintained to track changes to the table.

# Data Lake vs Delta Lake

	Data Lake	Delta Lake
<b>Transaction &amp; Isolation</b>	Multiple data pipelines reading and writing concurrently, great for scalability, but bad for ensuring data integrity	Enables ACID transactions, even serializable isolation level.
<b>Metadata handling</b>	No distributed processing of metadata	Treats metadata like regular data, utilizing distributed processing to handle metadata. Great for data that's truly big(PB scale, billions of files)
<b>Version control</b>	No version control	Provides data versioning
<b>Storage Format</b>	Different file formats, easy to dump data into, but hard to use	All data stored in Delta format
<b>Streaming Data</b>	Poor support for streaming data	Out-of-the-box support for streaming data
<b>Schema</b>	Hard to specify/enforce/change schema	Easy to specify/enforce/change schema due to unified file format

# Working with Delta

- **Tables** are equivalent to **dataframes** except
  - Table are defined at the workspace level, persists between notebooks
  - Data frames are defined at the notebook level
- Tables in Delta are generally classified in three levels
  - Raw
    - Bronze -> single source of truth
      - Raw, unprocessed data, with some metadata added
  - Silver
    - Cleaned, preprocessed data that is directly queryable
  - Gold
    - Highly refined views of data
      - Aggregated data
      - Feature tables



# Working with Delta

- Delta lakes are easy to create
  - When writing, simply specify “delta” instead of “parquet” or “csv”

```
data.write.format("delta").mode("overwrite").save("/tmp/delta-table")
```

- We can use Spark SQL directly on a directory of Delta data by specifying the directory path

```
SELECT * FROM delta.`/path/to/delta_directory`
```

- We can also create a table using Spark SQL

```
spark.sql("""DROP TABLE IF EXISTS customer_data_delta""")
spark.sql("""
  CREATE TABLE customer_data_delta
  USING DELTA
  LOCATION '{} '
  """.format(DataPath))
```

\* Since schema is already stored in metadata, we don't need to specify it when creating table

# Working with Delta - Append

- To append to delta lake, all we need to do is to change the mode of the previous write query

```
(newDataDF
  .write
  .format("delta")
  .partitionBy("Country")
  .mode("append")
  .save(DataPath)
)
```

- Changes in the data file will be reflected in the tables based on the file immediately

# Working with Delta - Upsert

- Upsert is used to simultaneously update and insert data

```
%sql
MERGE INTO customer_data_delta
USING upsert_data
ON customer_data_delta.InvoiceNo = upsert_data.InvoiceNo
  AND customer_data_delta.StockCode = upsert_data.StockCode
WHEN MATCHED THEN
  UPDATE SET *
WHEN NOT MATCHED
  THEN INSERT *
```

# Optimizations in Delta Lake

- Optimize
  - Performs **file compaction**
  - Small files are compacted together into new larger files **up to 1GB**
    - The 1GB size was determined by the Databricks optimization team as a trade-off between query speed and run-time performance when running Optimize.
  - Recommend to run Optimize on a daily basis, during off-hours, and increase/decrease frequency **based on business needs.**

# Optimizations in Delta Lake

- Partitioning
  - Used to speed up queries with **WHERE clause**
  - Relies on data being correctly partitioned
  - Will skip partitions that doesn't satisfy the WHERE condition
- ZOrdering
  - Another way to try to reduce the number of files searched for a query
  - Colocate related information in the same set of files.
  - Effectiveness goes down as more columns are considered.
- You can not use partitioning and ZOrdering on the same column

```
%sql  
OPTIMIZE delta_data_source  
ZORDER BY COL1,COL2...|
```

# Optimizations in Delta Lake

- Vacuum
    - Used to save some storage space
    - Cleans up invalid data files
      - Created by optimize/updates/upserts/deletions
- `VACUUM name-of-table RETAIN number-of HOURS;`
- It's recommended to **NOT** clean things younger than 7 days
    - old snapshots and uncommitted files can still be in use by concurrent readers or writers to the table.

# Using Time Travel in Delta Lake

- Viewing the history of a Delta table
  - This will give some metadata of all versions of this table
  - Result is in a tabular format with “version” (integer) being one of the columns

```
Cmd 13
1  %sql
2  DESCRIBE HISTORY table_name
```

- To query an older version of a Delta table, use **AS OF**
  - The AS OF clause should directly follow FROM, before WHERE conditions

```
Cmd 15
1  %sql
2  SELECT COUNT(*)
3  FROM table_name
4  VERSION AS OF 1
```